NVIDIA NEMOTRON 3 · NEBIUS TOKEN FACTORY

Your standing work, done every morning by a crew that asks before it acts.

A visual personal operations room where Ultra plans, Super reads and writes, Nano checks and learns, and every outgoing action waits for your approval.

Morning 1 (58,950 tokens)

$0.0216

24% saved by the cost router

Morning 2 (learned rule active)

$0.0209

27% saved · 16 memories learned

Replies sent without your OK

0

Hard approval gate on every side effect

THREE WALLS OF ALWAYS-ON AGENTS

Chat bubbles hide the work. Flagship models burn the budget.

Plenty of agents can read an inbox on a schedule. The real failure modes happen when you leave them running every day across your inbox, codebase, and research.

Wall 01 · Unit Economics

16× Overpaying

Routing every file read, retry loop, verification check, and memory extraction through one flagship model makes daily background chores too expensive to keep on.

Solved by 3-Tier Nemotron Routing

Wall 02 · Blind Execution

40 Hidden Calls

OpenClaw and Hermes give you a chat window, but a single bubble cannot show which agent read which file, why a command retried, or what each step cost.

Solved by Live Multi-Agent Timeline

Wall 03 · Static Prompts

Repeat Questions

Agents either act without asking or ask the exact same permission every morning because yesterday's approval never compiles into a revocable rule in `SKILL.md`.

Solved by Policy-Compiling Gates

VISUAL MULTI-AGENT OPERATIONS ROOM

See who did what, which model ran it, and what it cost.

Instead of a spinner in a chat box, Nimbus renders every agent on a live timeline. Click any segment to open the **Loop Inspector** (Gather, Act, Verify) and read the exact email, calendar invite, or test output that agent saw.

`Nemotron-3-Ultra-550b-a55b` $1.00 / $3.00 per 1M

Orchestrator · Plans the run, coordinates handoffs, gates side effects

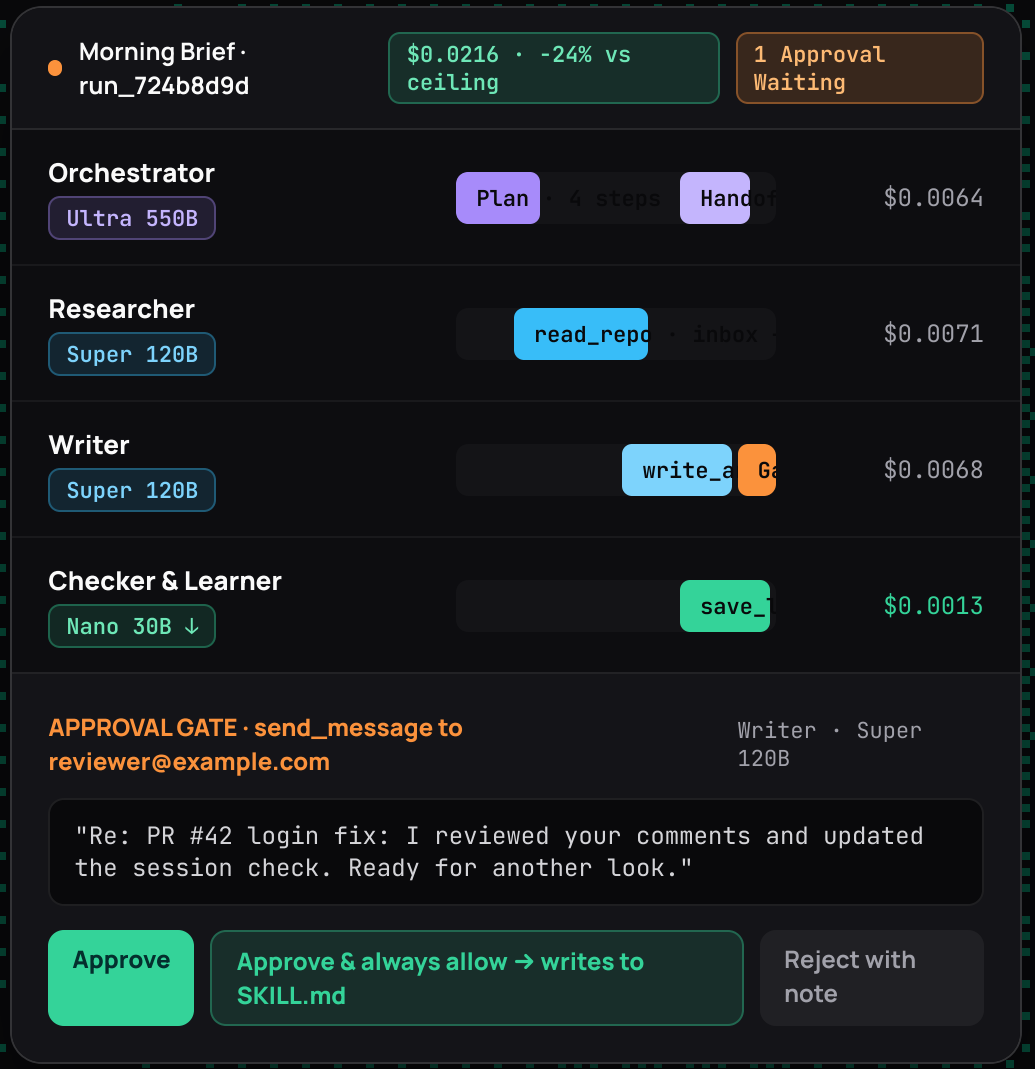
`nemotron-3-super-120b-a12b` $0.30 / $0.90 per 1M

Researcher, Analyst & Writer · 256K context for search, code & drafts

`NVIDIA-Nemotron-3-Nano-30B-A3B` $0.06 / $0.24 (1/16th Ultra)

Checker, Learner & Router Classifier · Fast verification & memory

One Row Per Agent · Model Badge & Live Price on Every Hop



RUN 2 STARTS WHERE RUN 1 LEFT OFF

Nothing leaves without you. Approved actions become rules you can revoke.

workstreams/ws-brief/SKILL.md · Rewritten by Nano Learner Revert skill change

```
# Morning Brief Skill & Learned Rules
+ allow: send_message to reviewer@example.com
+ source: fixtures/inbox/*.eml, fixtures/calendar.ics
## Learned Memories (16 active, top-k injected)
• Rank client contract deadlines above internal sync notes.
• Call send_message once per drafted reply so the approval card appears.
• Ground PR status strictly in github-notifications.json ("in review" vs "merged").
```

Learner output with `enable_thinking: false`: **102 tokens** Run 2 sends reviewer reply with 0 cards

Ship Nimbus Workstream

Self-Maintaining Repo Crew

Runs `pnpm eval` and `pnpm typecheck` in a sandbox, reproduces failing router goldens, writes a unified diff patch, and pauses for your OK before opening a GitHub PR.

OpenClaw & Hermes Bridge

POST /v1/chat/completions

Point OpenClaw or Hermes at Nimbus in 10 lines of config. Ask `"what needs me?"` and approve pending cards from Telegram or Slack while keeping one unified web timeline.

Every rule and memory can be forgotten or revoked in one click

Six live recording attempts ($0.49 total) exposed 7 real agent lessons.

01 · Nano 30B Thinking Budget

690 → 102 tok

With `<think>` on, Nano 30B spent its whole 1,500-token ceiling thinking and never called `save_learning`. Setting `enable_thinking: false` made it 6× smaller and 100% reliable.

Commit 23f32be

02 · Turn Budget & Synthesis

12 Turns + Report

The Researcher read one file per turn and exited at turn 8 with no report. Raising the budget to 12 with a tool-free final turn guaranteed a synthesis every run.

Commit be3ddc0

03 · Automatic Bus Handoffs

0 Re-reads

Ultra wrote `"(include the report)"` instead of pasting it, forcing the Writer to re-read every raw file. The runtime bus now attaches prior crew reports automatically.

Commit 3e117e6

04 · Rule Key Normalization

100% Policy Match

Normalizing `{ to, body }` and `{ channel, text }` across roles ensured saved `allow:` rules in `SKILL.md` matched on Run 2 with zero friction.

Verified in take 1915

ZERO MOCK FOOTAGE · 100% MEASURED

Try it in two minutes with zero keys, or live on Nebius.

```
git clone github.com/harrymove-ctrl/nimbus
&& pnpm dev
```

One morning, the first run

$0.0216

58,950 tokens · 23 model calls

The next morning (Run 2)

$0.0209

3% cheaper · starts from what Run 1 learned

Saved by the 3-tier router

24% → 27%

vs. pinning every hop to its ceiling model

A/B Trajectory Eval Suite

19 / 19

145 per-hop assertions · −47.9% router spend

Live Chat API Run (/v1)

$0.0282

Started & approved via /v1/chat/completions

Replies sent without your OK

0

Every side effect waits for approval or rule