NVIDIA NEMOTRON 3 · NEBIUS TOKEN FACTORY

# Your standing work, done every morning by a crew that asks before it acts.

A visual personal operations room where NVIDIA Nemotron 3 Ultra plans, Super reads and writes, Nano checks and learns, and every outgoing action waits for your approval.

Morning 1 (58,950 tokens)

$0.0216

24% saved by the 3-tier cost router

Morning 2 (learned rule active)

$0.0209

27% saved · 16 memories learned

Replies sent without your OK

0

Hard approval gate on every side effect

THREE WALLS OF ALWAYS-ON AGENTS

# Chat bubbles hide the work. Flagship models burn the budget.

Plenty of agents can read an inbox on a schedule. The real walls appear when you leave personal agents running every day.

Wall 01 · Unit Economics

## 16× Overpaying

Routing every file read, retry, verification check, and memory write through one flagship model makes daily chores too expensive.

Solved by 3-tier Nemotron routing



Wall 02 · Blind Execution

## 40 Hidden Calls

OpenClaw and Hermes give you a chat thread, not an operations room. You cannot see which agent read which file, why a command retried, or what each hop cost.

Solved by the Live Timeline

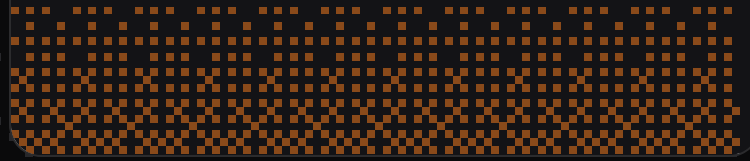


Wall 03 · Static Prompts

## Repeat Questions

Most assistants ask for the exact same permission every morning because yesterday's approval never updates the agent's skill file.

Solved by Policy-Compiling Gates



VISUAL MULTI-AGENT OPERATIONS ROOM

# See who did what, which model ran it, and what it cost.

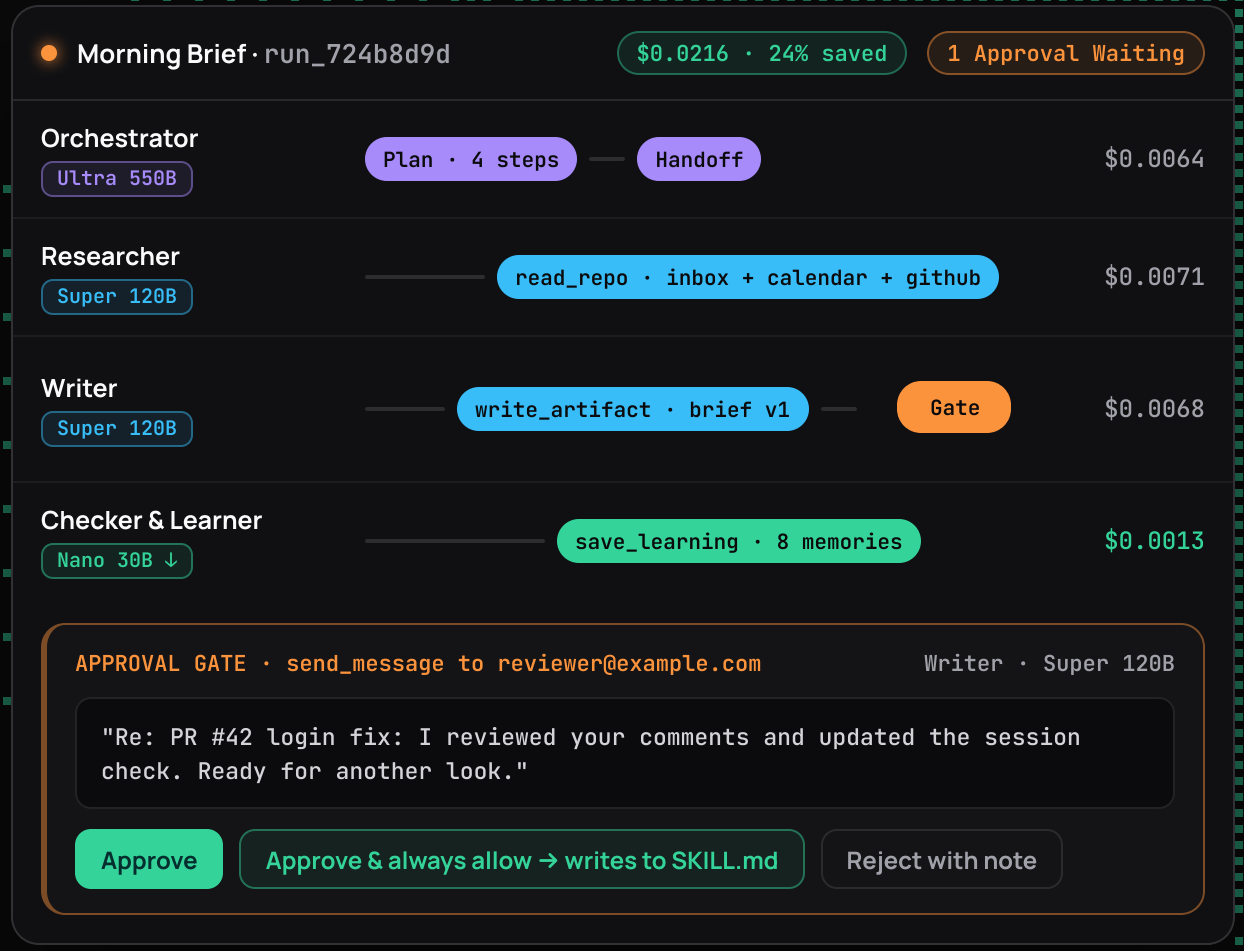
Instead of a spinner in a chat box, Nimbus renders every agent on a live timeline. Click any segment to open the Loop Inspector and read the exact email, calendar invite, or test output that agent saw.

**nvidia/Nemotron-3-Ultra-550b-a55b**
**$1.00 / $3.00** per 1M in / out
Orchestrator · Plans the run, coordinates handoffs, gates side effects

**nvidia/nemotron-3-super-120b-a12b**
**$0.30 / $0.90**
Researcher, Analyst & Writer · 256K context for inbox, code & drafts

**nvidia/NVIDIA-Nemotron-3-Nano-30B-A3B**
**$0.06 / $0.24** 1/16th the input cost of Ultra
Checker, Learner & Router Classifier

One row per agent · model badge and live price on every hop



RUN 2 STARTS WHERE RUN 1 LEFT OFF

# Nothing leaves without you. Approved actions become rules you can revoke.

workstreams/ws-brief/SKILL.md · Rewritten by Nano Learner

Revert skill change

```
# Morning Brief skill and learned rules
+ allow: send_message to reviewer@example.com
+ source: fixtures/inbox/*.eml, fixtures/calendar.ics

## Learned Memories (16 active, top-k injected via BAAI/bge-en-icl)
• Rank client contract deadlines above internal sync notes.
• Call send_message once per drafted reply so the approval card appears.
• Ground PR status strictly in github-notifications.json ("in review" vs "merged").
```

Learner output with `enable_thinking: false`: **102 tokens**

Run 2 sends the reviewer reply with 0 cards

Ship Nimbus Workstream (ws-ship)

## Self-Fixing Repo

Runs pnpm eval and pnpm typecheck in a sandbox, reproduces failing router goldens, writes a unified diff patch, and waits for your OK before opening a GitHub PR.

OpenClaw & Hermes Bridge

## POST /v1/chat

Point OpenClaw or Hermes at Nimbus in 10 lines of config. Ask "what needs me?" and approve pending cards from Telegram or Slack while keeping one web timeline.

WHAT REAL MODELS TAUGHT US BEYOND THE MOCK

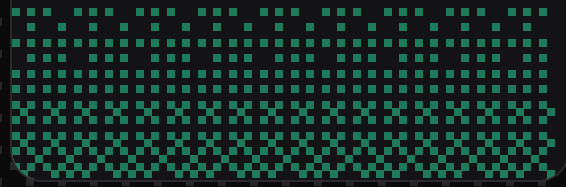
# 198 live Nebius Token Factory calls ($0.49 total) exposed 7 real agent lessons.

01 · Nano 30B Thinking Toggle

## 690 → 102 tok

With ‹think› on, Nano 30B spent its whole 1,500-token ceiling thinking and never called save_learning. Setting enable_thinking: false made it 6× smaller and 100% reliable.

Commit 23f32be

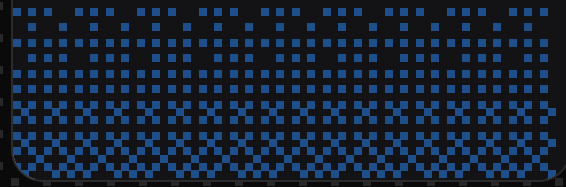


02 · Turn Budget & Synthesis

## 12 Turns + Report

The Researcher read one file per turn and exited at turn 8 with no report. Raising the budget to 12 with a tool-free final synthesis turn guaranteed a summary every run.

Commit be3ddc0

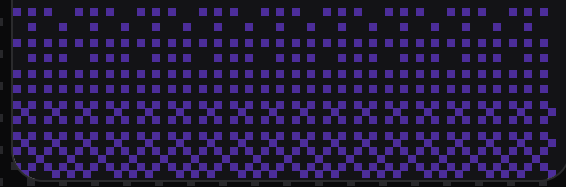


03 · Automatic Bus Handoffs

## 0 File Re-reads

Ultra wrote "(include the report)" instead of pasting it, forcing the Writer to re-read every raw file. The runtime bus now attaches prior crew reports automatically.

Commit 3e117e6



04 · Rule Key Normalization

## 100% Rule Match

Normalizing { to, body } and { channel, text } across roles ensured saved allow: rules in SKILL.md matched automatically on Run 2 with zero friction.

Verified in take 1915



ZERO MOCK FOOTAGE · 100% MEASURED

# Try it in two minutes with zero keys, or live on Nebius Token Factory.

```
git clone https://github.com/harrymove-ctrl/nimbus && cd nimbus && pnpm install && pnpm dev
```

One morning, the first run

**$0.0216**

58,950 tokens · 23 model calls (run_724b8d9d)

The next morning (Run 2)

**$0.0209**

3% cheaper · starts from what Run 1 learned (run_915dfc1d)

Saved by the 3-tier router

**24% → 27%**

vs. pinning every hop to its role's ceiling model

A/B Trajectory Eval Suite

**19 / 19**

145 per-hop assertions passing · 47.9% routed spend reduction

Live Chat API Run (/v1)

**$0.0282**

Started and approved via /v1/chat/completions (run_6b393ab6)

Replies sent without your OK

**0**

Every side effect waits for human approval or an explicit rule